

Ashlie Mulnix

Independent AI Safety Researcher

Selected work

- Built Solace-Vera, a 4-phase pre-action auditing pipeline (GitHub link below)
- Tested GPT, Gemini, Claude, Grok on 50 risk scenarios, two runs each
- Discovered that justifications change 100% of the time when risk scores are identical
- Published definition of MOC Effect on Zenodo
- Reported initial finding (MOC Effect) to OpenAI (out of scope). At a UC Berkeley researcher's suggestion, extended tests to Gemini, Claude, and Grok. Gained cross-model insights.

Background

- Self-taught, no CS degree
- Small business employee (florist)
- Two-year degree in Anthropology
- Two-year degree in Social Sciences

Links

- GitHub: <https://github.com/anchor-cloud/solace-vera-observability>
- Zenodo: <https://zenodo.org/records/19957469>
- Example evidence: Comparison of four models on a single high-risk scenario
https://github.com/anchor-cloud/solace-vera-observability/blob/main/blog_figures/fig_model_comparison_moc013.png

What I am looking for

Guidance on where to publish these findings, how to find collaborators to turn the scenario set into a benchmark, and whether to apply for a fellowship or research role.